

Position: Where Byte-Level Checkpoint Deduplication Dies: A Boundary Study of Semantic Redundancy in AI Training Checkpoints

Anonymous Author(s)

Abstract

Semantic checkpoint-deduplication systems report 6–896× storage savings by reasoning about *tensors*—identifying replicated, architecturally shared, or slowly changing parameters across ranks, jobs, and training steps. All of them require framework cooperation: offline tensor analysis, specialised grouped-I/O APIs, or accelerator-offloaded hashing. It is tempting to assume that an unmodified POSIX filesystem could capture the same savings with transparent, content-defined block deduplication. We test this assumption with one fixed measuring rod—BLAKE2b-128 fingerprints over fixed-size blocks—applied to real (CPU) GPT-2 small training runs in three checkpoint formats, precision and save-gap controls, synthetic multi-rank and multi-job fixtures, and, as positive controls, container image layers and multi-version source archives. The boundary is sharp: the byte layer sees the **copy axis** (DDP rank replicas, shared LoRA base files) at 100% hit ratio, but sees essentially **nothing** on the **similarity axis**—cross-step temporal redundancy ($\leq 0.29\%$ of 1 MB chunks, all zero blocks), same-architecture different-seed jobs ($\sim 0\%$), and disjoint parameter shards (0%)—even though a general compressor fares no better (zstd 1.08×; XOR-delta $\leq 1.28\times$). The same code path scores 59–98% on image layers and source trees. We argue the right role for a transparent layer is *learn-to-skip*: classify workloads before touching bytes, with asymmetric loss that can only ever affect performance, never correctness.

ACM Reference Format:

Anonymous Author(s). 2027. Position: Where Byte-Level Checkpoint Deduplication Dies: A Boundary Study of Semantic Redundancy in AI Training Checkpoints. In *Proceedings of 19th ACM Workshop on Hot Topics in Storage and File Systems (HotStorage'27)*. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Checkpoints already dominate AI-cluster storage traffic, and the problem grows with model size and checkpoint

frequency. Two FAST'26 systems attack the waste semantically: AdaCheck [6] classifies checkpoint tensors as replicated, architecturally shared, or temporally similar, and reports 6–896× savings; AITURBO [3] batches tensor I/O into grouped operations and computes BLAKE3 [8] fingerprints with optimised XPU kernels because CPU hashing cannot keep up with the storage fabric. The common denominator is cooperation: the application must expose tensor identity or use a specialised API. The closest byte-level predecessor, Kaiser et al. [5], characterised block-deduplication potential in *HPC* application checkpoints in 2016 and found application-dependent but real recurring savings—raising the question whether modern AI-training checkpoints behave the same; our answer is that they do not.

Meanwhile, transparent block deduplication is a textbook POSIX-filesystem feature [10, 11]: fingerprint content blocks, suppress redundant writes, and leave the application untouched. If semantic checkpoint redundancy were visible at the byte layer, *unmodified* PyTorch jobs would benefit simply by mounting the filesystem—no framework changes, no XPU. We asked exactly one question:

Of the redundancy that tensor-level systems exploit, how much reaches an unmodified POSIX byte layer?

We answer with a controlled measurement on a self-contained, reproducible GPT-2 small (124 M parameter) training fixture—real forward/backward and AdamW updates, not synthesised tensors—plus axis-specific fixtures and two workload families where byte deduplication is *expected* to win.

Findings. Cross-step identical 1 MB blocks are 0.14–0.28% for full checkpoints and **0%** for weights-only files; every hit is a zero block. Alignment sweeps, bf16 precision, and long save gaps do not help; neither zstd nor XOR-delta exploits the data. Yet identical DDP rank replicas and shared LoRA bases match at 100%, and image layers/source archives score 59–98% under the same code. Redundancy splits by **axis**, not by workload.

Contributions.

- (1) A controlled, reproducible study of redundancy visibility at the POSIX byte layer across three formats, two precisions, three save gaps, and all three redundancy axes.
- (2) A mechanism explanation—dense gradients plus decoupled weight decay make every fp32 byte change every step—with controls ruling out chunking, alignment, precision, and temporal distance.
- (3) A sharp copy/similarity boundary and its design implication: transparent storage should *learn to skip* hashing under an asymmetric-loss contract that keeps ML out of the correctness path.

2 Background and Redundancy Taxonomy

Checkpoint layouts. `torch.save` serialises a pickle object inside an uncompressed (STORED) zip container; tensor storage lands in numbered `data/N` entries. `torch.distributed.checkpoint` (DCP) writes one or more `.distcp` shard files plus metadata in the same zip-based format. `safetensors` [4] instead stores tensors as one contiguous little-endian byte region preceded by a JSON header. None of these formats compresses fp32 state, because it does not compress (§4).

Byte-level deduplication. A transparent client fingerprints fixed-size blocks (we use BLAKE2b-128) and suppresses blocks whose fingerprints already exist. Cost is dominated by hashing—why AITURBO computes BLAKE3 [8] on XPU kernels [3]—and by alignment: fixed grids miss offset-drifted content, which content-defined chunking [7, 12] fixes at extra CPU cost. Block stores from Venti [10] to iDedup [11] follow this model; iDedup also documented fragmentation cost on reads.

Three redundancy axes. Semantic systems exploit three structurally different sources of repeated bytes:

The crucial distinction is **copies versus similarities**. A copy is the same tensor placed in several files; the bytes are identical. A similarity is *the same parameter* in two semantic versions—updated by an optimizer, or initialised from a different seed. Semantic tensor keys (name, shape, allocation identity) recognise the latter; byte fingerprints can only match exact content. Table 1 is the taxonomy; Table 3 (§4.1) fills in what the byte layer actually sees on each axis.

3 Methodology

Fixture (inter-step axis, C). We implement a self-contained GPT-2 small (50,257 vocab, 1,024 positions, 768 width, 12 heads, 12 layers; tied input/output embeddings; pre-LayerNorm; HF-compatible tensor naming) and train it with AdamW (lr 10^{-3} ; betas (0.9, 0.999); eps

Table 1: Redundancy axes exploited by semantic checkpoint systems.

Axis	Example	Who benefits today
A. Intra-job copies	DDP replicates all parameters/optimizer state per rank; TP replicates embeddings	Any byte-identical copy mechanism
B. Inter-job similarity	Same architecture, different data/seed; LoRA jobs share one frozen base	Tensor-keyed systems [6]
C. Inter-step similarity	Consecutive checkpoints of one job	Tensor diff / grouped-I/O [3, 6]

10^{-8} ; weight decay 0.01) on random token streams (batch 1, sequence 128) for 20 steps, saving every step in three layouts: full `torch.save` (model+optimizer+step+loss; 1,493 MB), weights-only (498 MB), and DCP (1,494 MB)—about 65 GB total. A 100-step run provides long-gap and precision controls (torch 2.4.1+cpu, Python 3.8, seed 20260910; every script is checked in and re-runnable). We use a small CPU model because the mechanism under test (element-wise optimizer updates) is independent of parameter count.

Analyzer. For every checkpoint we fingerprint fixed blocks of {4 KB, 64 KB, 1 MB, 4 MB} using BLAKE2b-128, in a raw *file* view and a *zipdata* view where every `data/N` entry is chunked independently. We measure hit ratio against the previous version (*hit-prev*) and the union of all earlier versions (*hit-any*), all-pairs Jaccard, and offset-shift sweeps that re-chunk the reference after skipping {1, 16, 256, 4096} bytes—our detector for “identical content, drifted offset”, the condition CDC would rescue. A synthetic fixture of known random/zero/pattern segments validates exact instance counts.

Axis fixtures (A and B). From two independent 3-step runs we materialise four byte-identical full-checkpoint *rank replicas* per format (md5-verified), four disjoint 1/4 parameter *shards*, same-architecture cross-job pairs, and two LoRA-style jobs sharing one identical 498 MB base file with distinct 0.6 MB adapters.

Positive controls. Three images built sequentially on ubuntu:20.04, exported as flat rootfs tarballs and `docker save` OCI archives, plus `git archive` tarballs of five commits spanning ~276, 60, 60, and 20 commits.

Table 2: Steady-state (steps 5–19) cross-step hit-ratios. Every hit is a zero block.

Layout / view	4 KB	64 KB	1 MB	4 MB
<code>torch.save</code> / file	0.0037	0.0036	0.0014	0
<code>torch.save</code> / zipdata	0.0037	0.0037	0.0029	0
weights-only / file	0	0	0	0
DCP / file	0.0037	0.0036	0.0028	0

The image *application* layers are deliberately **full sibling layers**, not an optimised incremental chain. For archives we measure both a *per-file* view (each tar member chunked from offset zero) and the naive *raw* tar-stream view.

Scope. This is a controlled microbenchmark, not field-trace analysis; absolute ratios at hyperscale and non-AdamW optimisers are out of scope (§7).

4 Findings

4.1 Temporal redundancy is invisible at the byte layer

Every hit in the full-checkpoint rows (Table 2, Figure 1) is an all-zero block: exact-block ratio equals zero-block ratio (0.37%), and weights-only files score exactly zero at every size. At 1 MB, all-pairs Jaccard across steps 0–19 is a flat 0.000351 regardless of step distance. Temporal similarity—what AdaCheck and AITURBO harvest step after step—does not exist as equal bytes. Notably, AITURBO’s finest detection granularity is 4 MB [3]; at exactly that setting every format yields zero cross-step hits—XPU hashing would be spent on a stream containing nothing to find.

Not an alignment problem. Two complementary tests. *Within* checkpoints, shifting the previous checkpoint’s chunk grid by 1, 16, or 256 bytes leaves the hit ratio unchanged (e.g. 0.001404 at 1 MB for every shift; a 4096 shift removes the prefix chunk and nothing else). Identical-but-drifting content, the failure mode content-defined chunking fixes, is absent: there is no hidden redundancy for CDC to uncover (Figure 2).

Across workload families, alignment matters sharply. Chunking archive byte streams naively (raw tar view) collapses even our positive controls: rootfs 4 KB hits drop 58.6%→3.9% (v2) and 97.7%→10.4% (v3); source archives drop 94.8%→54.3% in the closest pair and to 1.4–4.5% at moderate distances; at 1 MB every raw-stream pair is zero, as tar reordering and headers destroy grid alignment. Checkpoint zipdata entries, in contrast, *are* chunked from independent aligned boundaries—and still score zero. Alignment is necessary (wrong boundaries read zero) but not sufficient.

Table 3: Axis map (per-file aligned blocks, same trained states).

Axis	Comparison	4 KB	64 KB	1 MB
A	DDP rank replicas	1.000	1.000	1.000
copies	(<code>torch.save</code> /safetensors)			
A	Disjoint FSDP/TP	0	0	0
divides	shards vs earlier shards			
B cross-	full ckpt / weights	0.000	0.000	0.0007
job	safetensors, diff. seed			
B	shared base file, job2 vs	1.000	1.000	1.000
LoRA	job1			
B	distinct adapter files	0	0	—
LoRA				
B	whole-job directories	0.9988	0.9988	1.000
LoRA				
C tem-	consecutive steps (Ta-	0.0037*	0.0036*	0–0.0029
poral	ble 2)			

*hits are zero blocks. The single 1 MB inter-job hit (1 of 1,423 blocks) is zero content, not structure.

Not precision, age, or compressibility. *Precision control.* bf16 weights across 100 steps yield 6–11 matching 4 KB blocks out of 60,726 ($\leq 0.018\%$) and zero at 64 KB or 1 MB; fp32 weights are exactly zero at gaps of 1, 50, and 100 steps. Lower mantissa precision does not quantise Adam updates into byte-identical blocks.

Format control. A safetensors run replaying the **identical** 20-step sequence (bit-identical losses: 10.989, 10.916, 10.968, ...) scores 0.0025% hits at 4 KB and zero above—the result is a property of the bytes, not the zip container.

Compression control. Perhaps redundancy exists in a form blocks cannot express (Figure 3). zstd-1 on one checkpoint compresses 1.08×; an XOR-delta stream between adjacent steps compresses only 1.15–1.28×; byte-wise equality between adjacent steps is just 19.8–30.3%. fp32 training state in motion carries almost no extractable redundancy by any measure.

Mechanism: every byte is touched every step. The zero result has a precise mechanism. (1) AdamW applies decoupled weight decay: every participating element receives $p \leftarrow p - lr \cdot wd \cdot p$ each step, a relative change of 10^{-5} —small but guaranteed to flip fp32 low bytes. (2) Adam momentum slots decay and accumulate per element. (3) In our tied-head GPT-2, the lm-head backward produces a dense gradient through the whole embedding matrix, and the embedding backward materialises **zero (not None) gradients** for unsampled rows, so weight decay touches those rows too—we verified that rows unused by the current batch still change bytes every step. At fp32, essentially every parameter byte changes each step, while tensor *keys* (names, shapes, identity) stay constant: the whole axis boundary in one sentence. **Redundancy-axis map: copies visible, similarities invisible.**




Figure 1: Cross-step hit-any ratio per checkpoint format and chunk size (log scale). The byte layer finds only zero blocks on training checkpoints (red), versus 59–98% hits for container image layers and source archives.




Figure 2: Hit-any ratio vs. chunk size: no size recovers cross-step redundancy.




Figure 3: Adjacent-step byte equality is only 19.8–30.3%; zstd and XOR-delta zstd find almost no redundancy either.

Whole-object copies—rank replicas and shared base files—collapse to one physical copy transparently; the LoRA directory result (99.9%) shows a transparent store




Figure 4: Hit ratio vs. version distance (log x): checkpoints stay at the zero-block floor; source/rootfs overlaps stay high.

already handles the dominant bytes of the LoRA pattern, missing only the 0.6 MB adapters. Everything requiring recognition of “the same parameter in another version”—consecutive steps, independent jobs, disjoint shards—is invisible. The 6–896× headline savings of semantic systems live on this invisible side; their framework requirements are forced by where the redundancy is. **Positive controls: the same instrument scores 59–98%.**

The 75 MB ubuntu base blob is shared verbatim across all three images; even the 191/197 MB sibling application layers—full reinstalls—contain 59% and 98% of blocks seen earlier, because dpkg overwrites the same library and binary files with identical content (Table 4, Figure 4). System files are byte-stable across versions; optimizer state is byte-stable for exactly one step. With the same code scoring 98% and 0%, the zero cannot be blamed on the instrument.

Table 4: Positive controls (per-file view, hit-any).

Workload	4 KB	64 KB	1 MB
Image rootfs v2 vs v1	0.586	0.623	0.660
Image rootfs v3 vs v1+v2	0.977	0.974	0.981
Source, ~276 commits apart	0.324	0.115	—
Source, moderate gaps	0.705/0.754	0.563/0.585	—
Source, ~20 commits apart	0.948	0.956	—
Training checkpoints (all formats)	0	0	0

5 Implications: Learn to Skip, Not to Deduplicate

A classifier before the hash. If hashing is expensive (AITURBO’s XPU is the admission [3]) and temporal checkpoints never hit, a transparent client should not be smarter about *how* it hashes checkpoint streams—it should decide *whether to hash at all*, from signals available before touching bytes: write pattern (periodic large sequential file from a training job), file size/age, directory topology (rank replicas), and cheap prefix probes. Copy-heavy paths (rank replicas, base models, image pulls, log/tar bundles) always hash; churn-heavy training state skips to write-through.

Asymmetric loss keeps ML off the correctness path. A gating classifier makes one mistake in each direction: a false positive (hashed a unique stream) wastes fingerprints; a false negative (skipped a redundant stream) loses one saving. Both affect performance only. Suppression still requires an exact fingerprint match, so ML can neither corrupt data nor lose a byte—unlike prediction-in-the-data-path systems such as LinnOS [2] or KML’s learned kernel heuristics [1]: prediction proposes, exact matching disposes. Our measurements supply the labelled **negative** workloads such a gate previously lacked; earlier write-side prediction only knew workloads where redundancy is the default.

A lower bound, not an obituary, for semantic systems. The zero on the similarity axis quantifies the *necessity* of tensor cooperation: no transparent cleverness—CDC, larger blocks, compression—can intercept savings whose identical units are tensor keys, not bytes. The 100% on the copy axis, conversely, shows transparent dedup remains right for replicated saves and shared bases; a deployed system should compose both.

Existence proof. In our research POSIX filesystem, an inline kernel content-hash fast path with a gated write needle cuts dedup metadata RPCs from 1,024 to 258 per 256MB write round and improves steady-state write bandwidth by 50% across four rounds, with

md5-identical output. The gated path is cheap and safe; this paper determines *where it must not be switched on*. (System name withheld for blind review.)

6 Related Work

Semantic checkpoint storage. AdaCheck [6] rewrites tensors across parallelism, architecture, and temporal dimensions; AITURBO [3] adds grouped I/O and XPU BLAKE3 fingerprinting. Kaiser et al. [5], the closest predecessor, measured byte-level deduplication potential across HPC checkpoints, where static memory regions produced application-dependent savings. We revisit the question for AI-training state under AdamW, splitting redundancy into copy and similarity axes: the answer flips for temporal similarity but holds for copies.

Block deduplication and chunking. Venti [10] established content-addressed blocks; iDedup [11] studied inline deduplication and its fragmentation cost; LBFS [7] introduced and FastCDC [12] optimised content-defined chunking. Our shift sweep tests CDC’s precondition (offset-drifted identical content) directly and finds it absent.

Learning in storage. LinnOS [2] reroutes I/Os from a learned latency model; KML [1] sets kernel heuristics such as readahead with learned models. In both, model output directly changes the I/O action. Our gate differs structurally: ML decides only whether an exact algorithm runs, never what data it returns.

AI storage workloads. Mooncake [9] accelerated KVCache-centric serving traffic, a separate track we do not enter. Traditional-workload baselines (FIU/MSR traces) are deferred to the full-paper version.

7 Limitations and Conclusion

Limitations. One architecture (124M GPT-2), one optimiser family (AdamW), random-token training, CPU, and single-process fixtures: absolute numbers will move at scale, but the mechanism (per-element updates flip fp32 bytes; similarity keys are tensor names) is architecture-independent, and our controls isolate it directly. Multi-node traces, non-AdamW optimisers, and a full CDC implementation (we give strong precondition evidence rather than an implementation) are clear extensions, as are read-side fragmentation effects on restore.

Conclusion. Redundancy splits sharply at the POSIX boundary: the byte layer sees every *copy* (100%) and nearly none of the *similarity* ($\approx 0\%$) semantic systems monetise, while the same code scores 59–98% on images and sources. Transparent storage cannot earn that 6–896 $\times$; it should learn to skip churn and reserve exact fingerprinting for where copies live.

References

- [1] Ibrahim Umit Akgun, Ali Selman Aydin, Ahmad Shaikh, Lukas Velikov, and Erez Zadok. 2021. A Machine Learning Framework to Improve Storage System Performance. In *13th ACM Workshop on Hot Topics in Storage and File Systems (HotStorage '21)*. ACM. doi:10.1145/3465332.3470875
- [2] Mingzhe Hao, Levent Toksoz, Nanqin Li, Edward Edberg Halim, Henry Hoffmann, and Haryadi S. Gunawi. 2020. Lin-nOS: Predictability on Unpredictable Flash Storage with a Lightweight Neural Engine. In *14th USENIX Symposium on Operating Systems Design and Implementation (OSDI '20)*. USENIX Association, 629–646.
- [3] Yingyi Hao, Ting Yao, Xingda Wei, Dingyan Zhang, Tianle Sun, Yiwen Zhang, Zhiyong Fu, Huatao Wu, and Rong Chen. 2026. Fast Cloud Storage for AI Jobs via Grouped I/O API with Transparent Read/Write Optimizations. In *24th USENIX Conference on File and Storage Technologies (FAST '26)*. USENIX Association. System name: AITURBO.
- [4] Hugging Face. 2024. safetensors Serialization Format. Software specification. <https://huggingface.co/docs/safetensors>
- [5] Jürgen Kaiser, Ramy Gad, Tim Süß, Federico Padua, Lars Nagel, and André Brinkmann. 2016. Deduplication Potential of HPC Applications' Checkpoints. In *2016 IEEE International Conference on Cluster Computing (CLUSTER '16)*. IEEE, 413–422. doi:10.1109/CLUSTER.2016.32
- [6] Weijie Liu, Shengwei Li, Zhiquan Lai, Keshi Ge, Qiaoling Chen, Peng Sun, Dongsheng Li, and Kai Lu. 2026. AdaCheck: An Adaptive Checkpointing System for Efficient LLM Training with Redundancy Utilization. In *24th USENIX Conference on File and Storage Technologies (FAST '26)*. USENIX Association, 271–289.
- [7] Athicha Muthitacharoen, Benjie Chen, and David Mazières. 2001. A Low-Bandwidth Network File System. In *Proceedings of the 18th ACM Symposium on Operating Systems Principles (SOSP '01)*. ACM, 174–187. doi:10.1145/502034.502052
- [8] Jack O'Connor, Jean-Philippe Aumasson, Samuel Neves, and Zooko Wilcox-O'Hearn. 2020. BLAKE3: One Function, Fast Everywhere. Specification and reference implementation. <https://blake3.io>
- [9] Ruoyu Qin, Zheming Li, Weiran He, Jiale Cui, Feng Ren, Mingxing Zhang, Yongwei Wu, Weimin Zheng, and Xinran Xu. 2025. Mooncake: Trading More Storage for Less Computation—A KVCache-centric Architecture for Serving LLM Chatbot. In *23rd USENIX Conference on File and Storage Technologies (FAST '25)*. USENIX Association, 155–170. Best Paper.
- [10] Sean Quinlan and Sean Dorward. 2002. Venti: A New Approach to Archival Storage. In *1st USENIX Conference on File and Storage Technologies (FAST '02)*. USENIX Association, 89–101.
- [11] Kiran Srinivasan, Tim Bisson, Garth R. Goodson, and Kaladhar Voruganti. 2012. iDedup: Latency-aware, Inline Data Deduplication for Primary Storage. In *10th USENIX Conference on File and Storage Technologies (FAST '12)*. USENIX Association.
- [12] Wen Xia, Yukun Zhou, Hong Jiang, Dan Feng, Yu Hua, Yuchong Hu, Qing Liu, and Yucheng Zhang. 2020. The Design of Fast Content-Defined Chunking for Data Deduplication Based Storage Systems. *IEEE Transactions on Parallel and Distributed Systems* 31, 9 (2020), 2017–2031. doi:10.1109/TPDS.2020.2984601